

Native Mode-Based Optimizations of Remote Memory Accesses in OpenSHMEM for Intel Xeon Phi

*Naveen Namashivayam, Sayan Ghosh, Dounia Khaldi, Deepak Eachempati and
Barbara Chapman*

Department of Computer Science
University of Houston

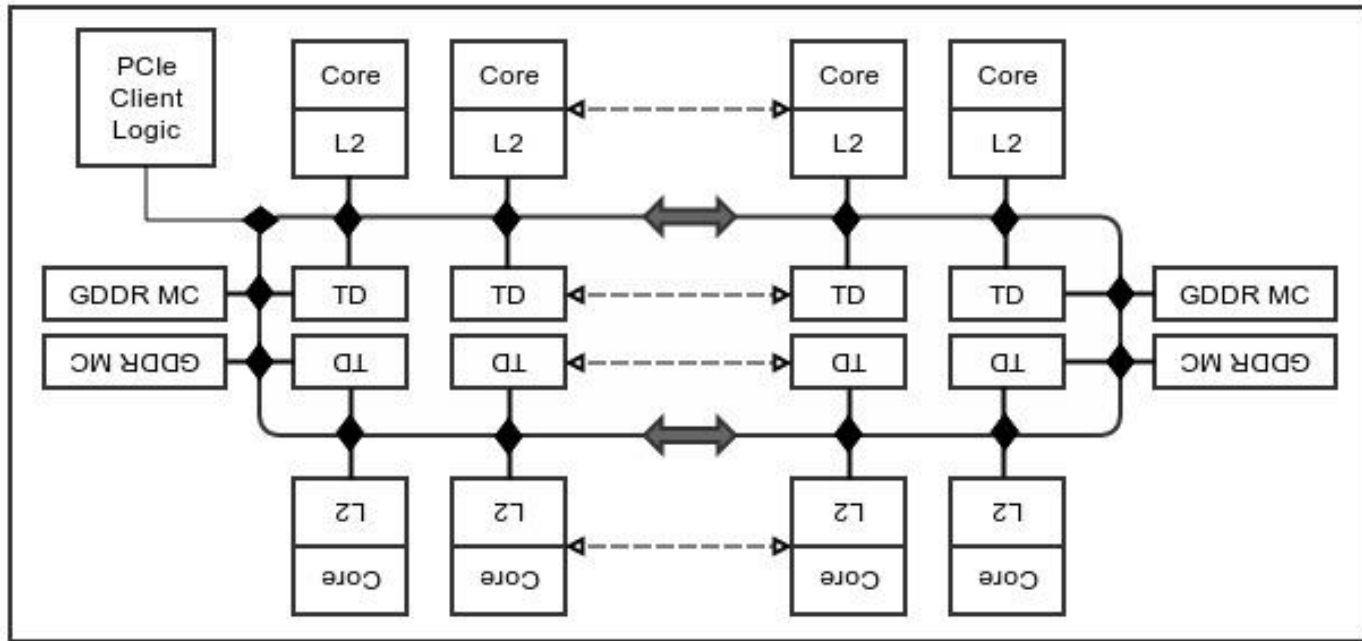
Motivation

- PGAS
 - An emerging programming model to be explored
 - OpenSHMEM on MIC in Native Mode
- Intel Xeon Phi: Many Integrated Core (MIC)
- Knight's Corner (KNC)
 - Current generation
 - Modified x86 co-processor
 - Unlike GPGPUs, porting applications is straightforward
- Knight's Landing (KNL)
 - Future Generation
 - Main CPU

Contents

- **Xeon Phi Architecture**
- Why OpenSHMEM?
- RMA put/get vs. Local load/store
- Optimizing Collectives: Reduction as a use case
- Experimental Results: NAS parallel benchmarks


Xeon Phi Architectural Overview




- $61 * 4 = 244$ cores
- Ring Interconnect
- Bidirectional
- x86 compatibility
- L1(32KB), L2(512KB)
- 8 Memory Controllers
- 16 GDDR5 memory
- PCI Express System

Xeon Phi – Modes of Programming


Native Mode



Offload Mode



Symmetric Mode



Contents

- Xeon Phi Architecture
- **Why OpenSHMEM?**
- RMA put/get vs. Local load/store
- Optimizing Collectives: Reduction as a use case
- Experimental Results: NAS parallel benchmarks

Why OpenSHMEM ?

- Bottleneck in selecting other PGAS models
 - Intel Compiler and MIC dependency
 - Fortran Coarray vs. OpenSHMEM
- Shared and distributed memory machines
- Reference implementation* developed by University of Houston is based on GASNet
- OpenSHMEM version for Intel Xeon Phi
- Test Bed -> Stampede Super Computer

* <http://openshmem.org/>

Contents

- Xeon Phi Architecture
- Why OpenSHMEM?
- **RMA put/get vs. Local load/store**
- Optimizing Collectives: Reduction as a use case
- Experimental Results: NAS parallel benchmarks

RMA put/get → Local load/store

- *shmem_ptr*: address of a data object on a specific PE
- No function call → enhancing compiler optimizations
- Enables optimizations such as vectorization


```
shmem_int_put(target, source, i, pe);
```



```
int *ptr = (int *)shmem_ptr(target, pe);  
for (m = 0; m < i; m+=1)  
    ptr[m] = source[m];
```



put/get → load/store - Latency

- Latency comparison
- *shmem_put/shmem_get* vs. load/store
- PGAS-Microbenchmarks from University of Houston



put/get → load/store - Bandwidth

- Bandwidth comparison
- *Blocksize* is represented in number of integers
- Results for 2 pairs and 64 pairs of PEs



Contents

- Xeon Phi Architecture
- Why OpenSHMEM?
- RMA put/get vs. Local load/store
- **Optimizing Collectives: Reduction as a use case**
- Experimental Results: NAS parallel benchmarks

Optimizing SHMEM Collectives

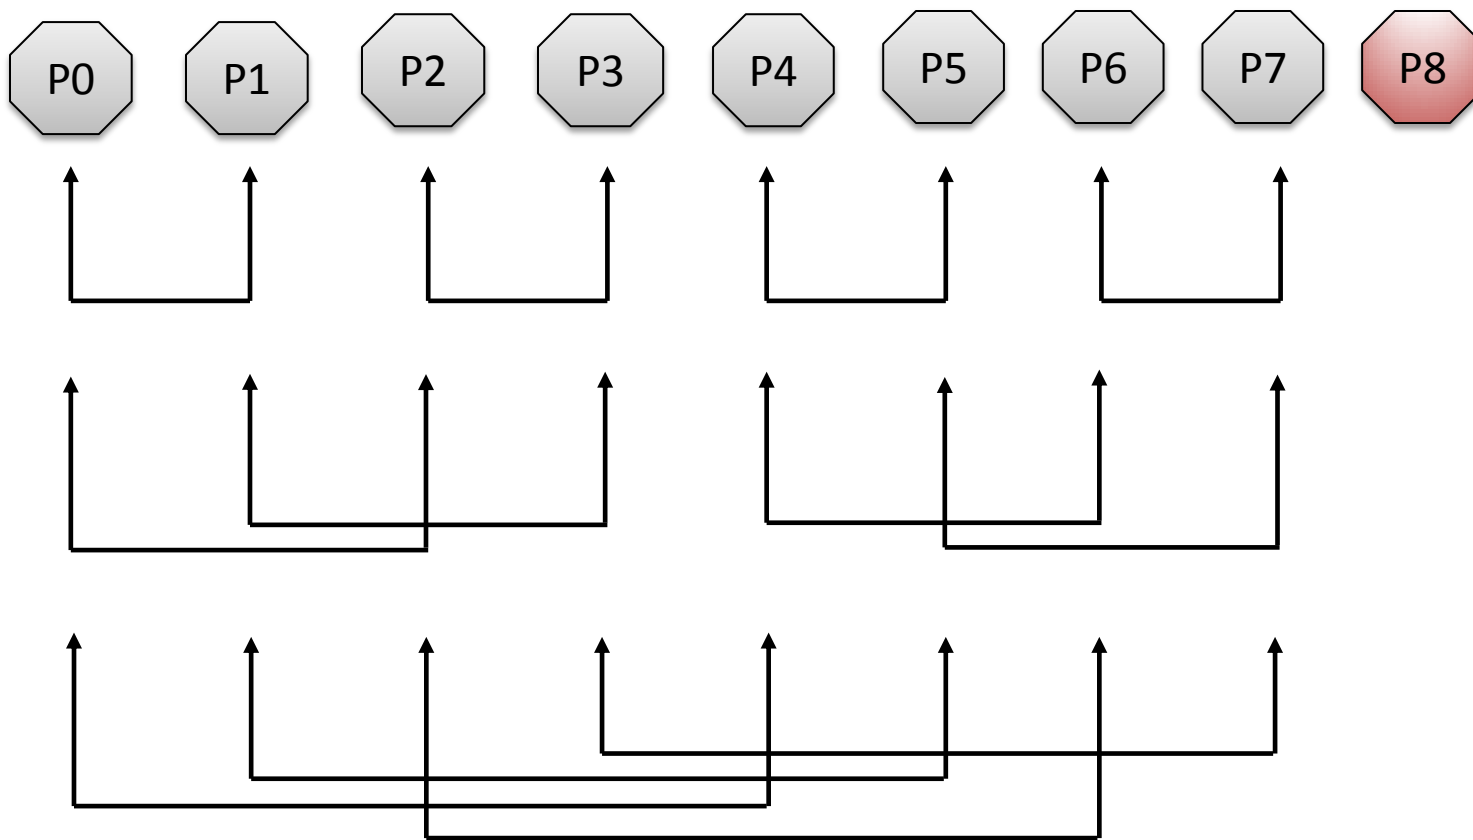
- Using load/store operations to optimize SHMEM collectives: Barrier, Broadcast, Collect and Reduce
- Implementation of different reduction algorithms:
 - Flat Tree (FT)
 - Recursive Doubling (RD)
 - Bruck (BR)
 - Rabenseifner: Reduce Scatter All Gather (RSAG)*
- Performance comparison with Intel MPI and MVAPICH
 - PGAS-Microbenchmarks from University of Houston[^]

* Rajeev Thakur, Rolf Rabenseifner, and William Gropp. Optimization of Collective Communication Operations in MPICH, 2005

[^] <https://github.com/uhhpctools/pgas-microbench>

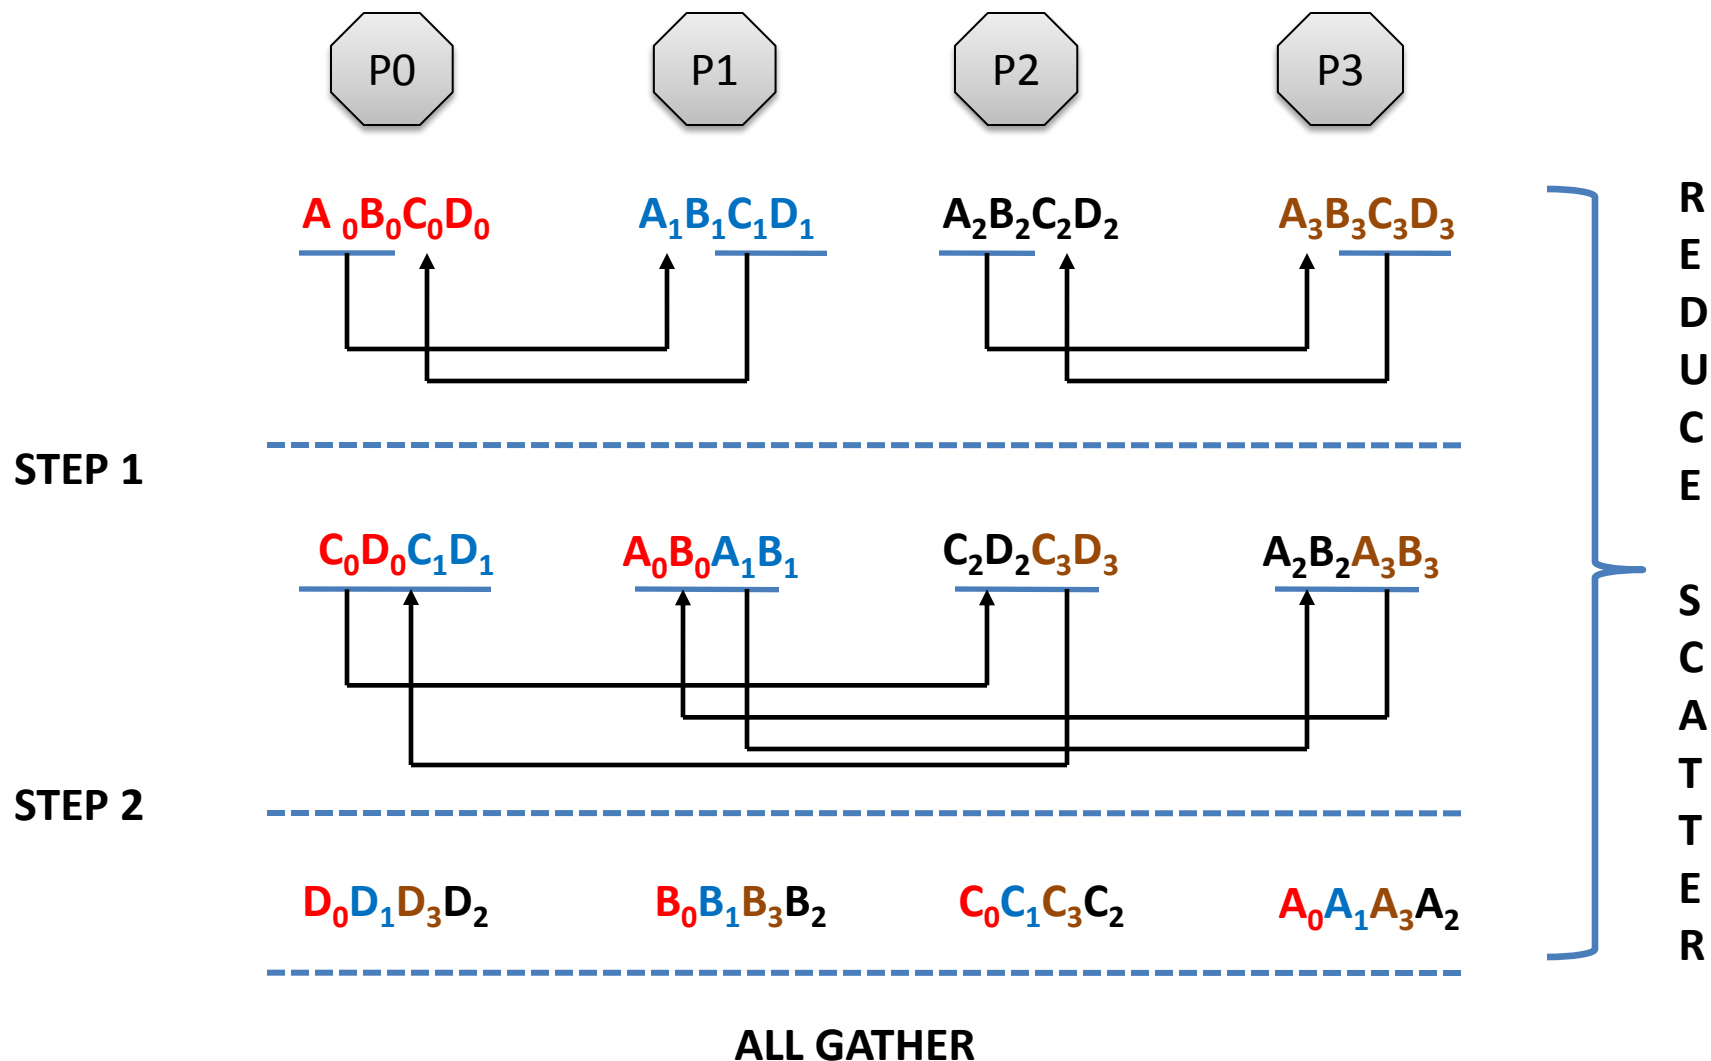
Recursive Doubling Algorithm

Array = { 1, 2 100 }



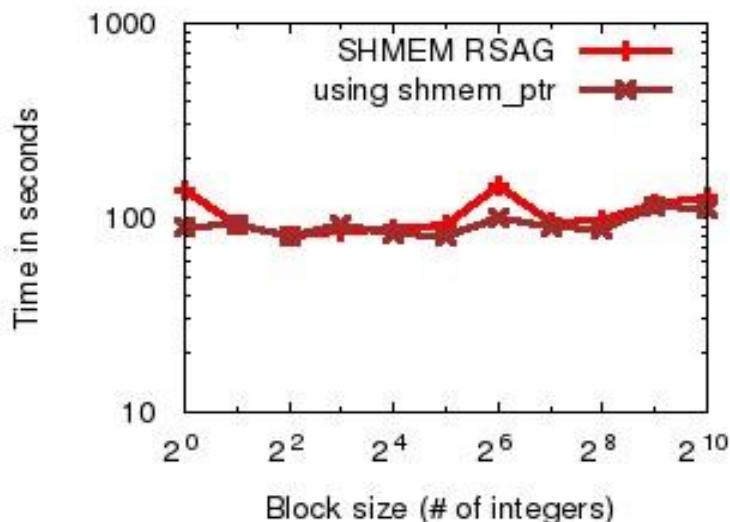
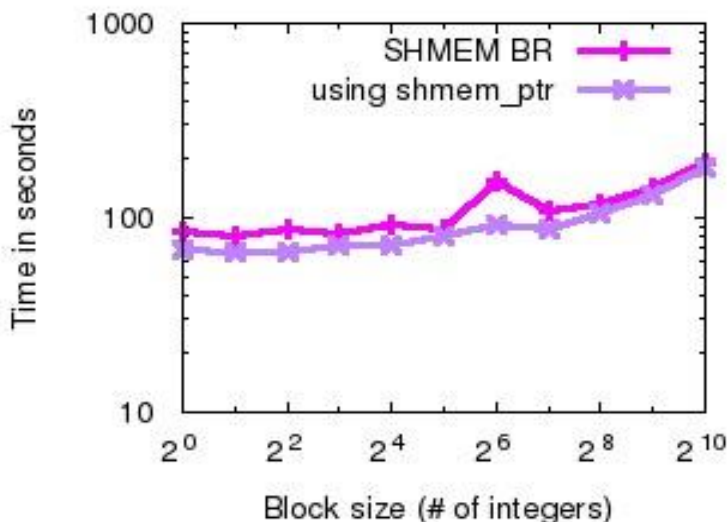


Reduce Scatter All Gather Algorithm

Array = { 1, 2 100 } \longrightarrow A = { 1, .. 25 }, B = { 26, .. 50 }, C = { 51, .. 75 }, D = { 76, .. 100 }



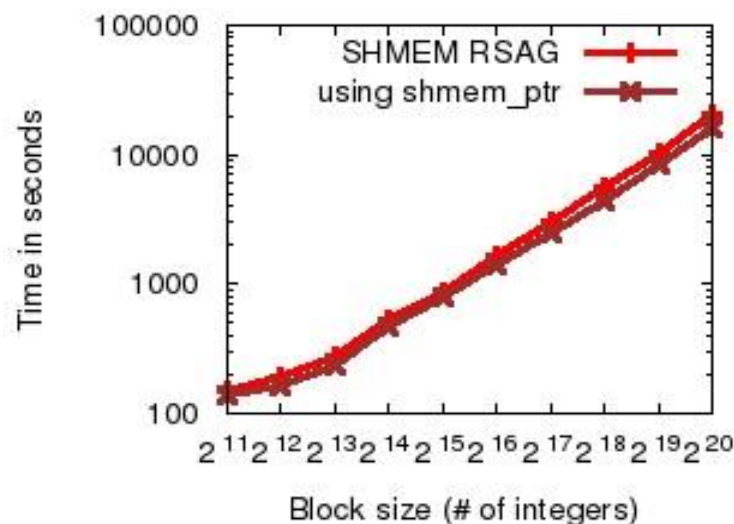

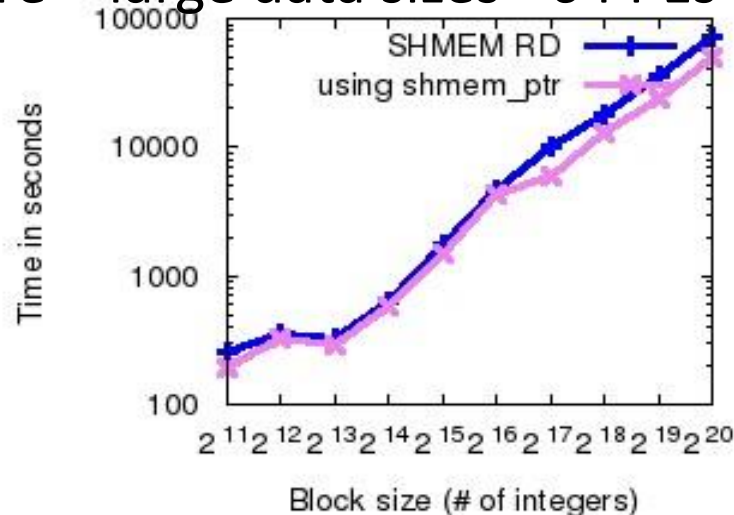
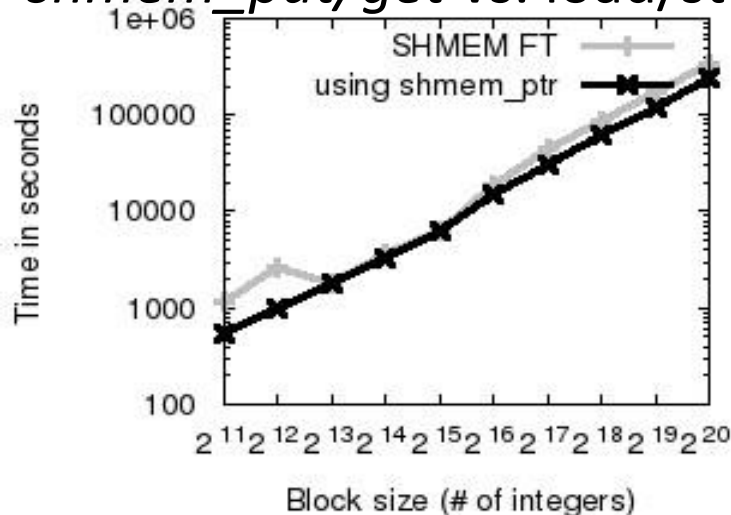
Reduction Algorithms Comparisons (1)

- shmem_put/get* vs. load/store – small data sizes - 64 PEs





Reduction Algorithms Comparisons (2)

- shmem_put/get* vs. load/store – large data sizes - 64 PEs



Reduction Algorithms Comparisons (3)

Various algorithms using load/store vs. Intel MPI and MVAPICH2 - 64 PEs



Contents

- Xeon Phi Architecture
- Why OpenSHMEM?
- RMA put/get vs. Local load/store
- Optimizing Collectives: Reduction as a use case
- **Experimental Results: NAS parallel benchmarks**

NAS Parallel Benchmarks


- MPI, CLASS C
- NAS SHMEM benchmarks*

Benchmark	Reduction (%)	Remote Access (%)	SHMEM* Ver. Available
MG	0.1	19.6	YES
BT	0	15.5	YES
EP	1.6	0	YES
SP	0	44.1	YES
IS	12.4	11.7	YES
CG	0	33.2	NO
FT	0.8	31.2	NO
DT	0	10	NO
LU	0.1	14.8	NO


* <https://github.com/openshmem-org/openshmem-npbs>

Experimenting on NAS Benchmarks (1)

IS NAS Benchmark – *shmem_put/get* vs. load/store




- Class B



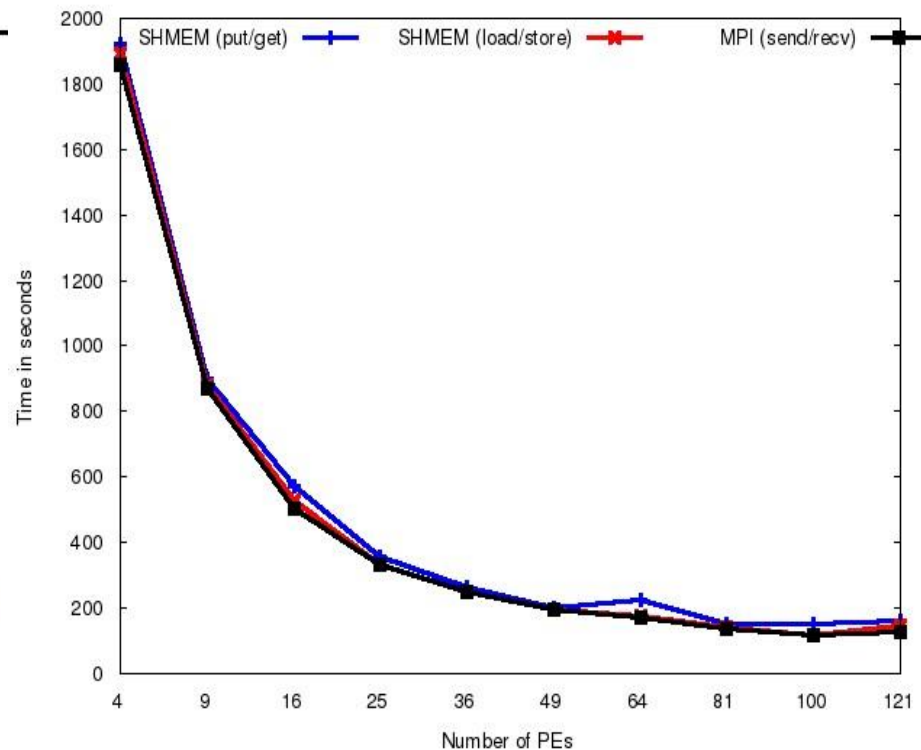
- Class C

Experimenting on NAS Benchmarks (2)

SP NAS Benchmark – *shmem_put/get* vs. load/store



- Class B

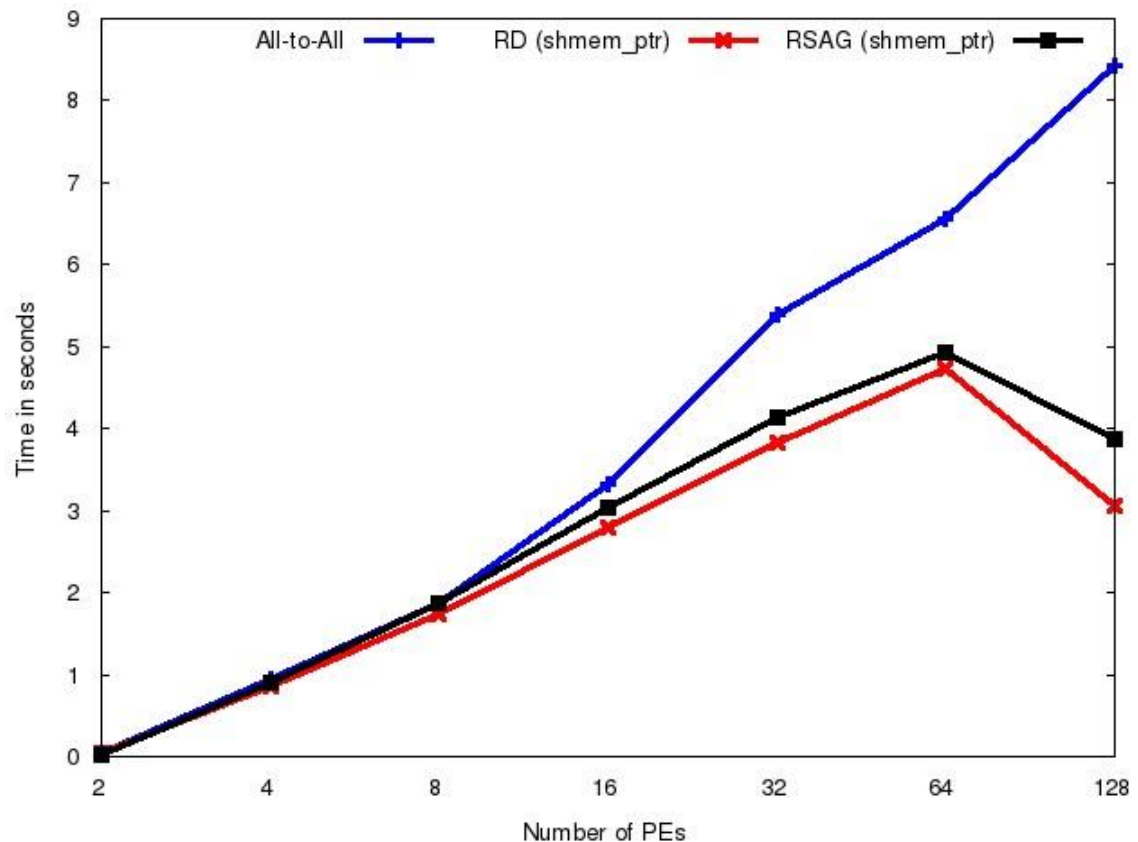


- Class C

Experimenting on NAS Benchmarks (3)

- IS NAS Benchmark, CLASS C
- load/store with various reduction algorithms
- All-to-All: Default reduction in

OpenSHMEM
reference
implementation



Conclusion

- Improved reduction algorithms: up to 22% compared to MVAPICH and 60% compared to IMPI
- Improved communications: decrease in latency by up to 60% and increase in bandwidth by up to 12x
- Use RD for small message sizes and RSAG for large message sizes.

Future Work

- Extending the reduction optimizations to other collectives such as Barriers
- Automation of translating RMA calls into load/store using OpenUH for shared memory systems
- PGAS Language-based on MIC, such as Fortran Coarray

Acknowledgments

- TOTAL
- TACC, Texas Advanced Computing Center



Native Mode-Based Optimizations of Remote Memory Accesses in OpenSHMEM for Intel Xeon Phi

*Naveen Namashivayam, Sayan Ghosh, Dounia Khaldi, Deepak Eachempati,
Barbara Chapman*

Department of Computer Science
University of Houston